

Applied Machine Learning

Regularization

Reihaneh Rabbany



Learning objectives

- intuition for model complexity and overfitting
- regularization penalty (L1 & L2)
- probabilistic interpretation

Linear regression

model:

$$\hat{y} = f_w(x) = w^\top x : \mathbb{R}^D \rightarrow \mathbb{R}$$

cost function:

$$J_w = \frac{1}{N} \sum_n \frac{1}{2} (y^{(n)} - \hat{y}^{(n)})^2 = \frac{1}{2} \|y - Xw\|^2$$

how to find w^* ?

closed form solution:

$$w^* = (X^\top X)^{-1} X^\top y$$

Or use
gradient
descent

partial derivatives:

$$\frac{\partial}{\partial w_d} J_w = \frac{1}{N} \sum_n (\hat{y}^{(n)} - y^{(n)}) x_d^{(n)}$$

gradient (all partial derivatives):

$$\nabla J(w) = \frac{1}{N} \sum_n (\hat{y}^{(n)} - y^{(n)}) x^{(n)} = \frac{1}{N} X^\top (\hat{y} - y)$$

repeat until stopping criterion:

$$w^{\{t+1\}} \leftarrow w^{\{t\}} - \alpha \nabla J(w^{\{t\}})$$

what if **linear fit is not the best**?

how to increase the model's expressiveness?

⇒ use nonlinear basis to create new nonlinear features from the existing ones

Nonlinear basis functions

replace original features in $f_w(x) = \sum_d w_d x_d$

with nonlinear bases $f_w(x) = \sum_d w_d \phi_d(x)$

linear least squares solution $(\Phi^\top \Phi)w^* = \Phi^\top y$

replacing X with Φ

a (nonlinear) feature


$$\Phi = \begin{bmatrix} \phi_1(x^{(1)}), & \phi_2(x^{(1)}), & \cdots, & \phi_D(x^{(1)}) \\ \phi_1(x^{(2)}), & \phi_2(x^{(2)}), & \cdots, & \phi_D(x^{(2)}) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_1(x^{(N)}), & \phi_2(x^{(N)}), & \cdots, & \phi_D(x^{(N)}) \end{bmatrix}$$

one instance

Nonlinear basis functions


examples

original input is scalar $x \in \mathbb{R}$




polynomial bases

$$\phi_k(x) = x^k$$



Gaussian bases


$$\phi_k(x) = e^{-\frac{(x-\mu_k)^2}{s^2}}$$




Sigmoid bases

$$\phi_k(x) = \frac{1}{1+e^{-\frac{x-\mu_k}{s}}}$$

Example: Gaussian bases



$$\phi_k(x) = e^{-\frac{(x-\mu_k)^2}{s^2}}$$



$$y^{(n)} = \sin(x^{(n)}) + \cos(\sqrt{|x^{(n)}|}) + \epsilon$$

prediction for a new instance


$$f(\mathbf{x}') = \boldsymbol{\phi}(\mathbf{x}')^\top (\Phi^\top \Phi)^{-1} \Phi^\top \mathbf{y}$$

new instance


\mathbf{w} found using LLS

features evaluated for the new point

Example: Gaussian bases




$$\phi_k(x) = e^{-\frac{(x-\mu_k)^2}{s^2}}$$




our fit to data using **10 Gaussian bases**

why not more?

Example: Gaussian bases




$$\phi_k(x) = e^{-\frac{(x-\mu_k)^2}{s^2}}$$




using 50 bases!

Example: Gaussian bases




$$\phi_k(x) = e^{-\frac{(x-\mu_k)^2}{s^2}}$$




using 200, thinner bases ($s=.1$)

Generalization?




$D = 5$



$D = 10$

$D = 50$




$D = 200$

lower training error →

which one of these models performs better at **test time**?


Overfitting

which one of these models performs better at **test time**?



An observation


when overfitting, we sometimes see large weights




dashed lines are $w_d \phi_d(x) \quad \forall d$

$$f_w(x) = \sum_d w_d \phi_d(x)$$


$D = 10$



$D = 17$



$D = 20$



idea: penalize large parameter values

Ridge regression

also known as

L2 regularized linear least squares regression:

$$J(w) = \frac{1}{2} \|Xw - y\|_2^2 + \frac{\lambda}{2} \|w\|_2^2$$

$\frac{1}{2} \sum_n (y^{(n)} - w^\top x)^2$

sum of squared error

$w^\top w = \sum_d w_d^2$

squared L2 norm of w

regularization parameter $\lambda > 0$ controls the strength of regularization

a good practice is to **not** penalize the intercept $\lambda(\|w\|_2^2 - w_0^2)$




λ is a hyper-parameter (use a validation set or cross-validation to pick the best value)

Ridge regression example

Visualizing the effect of regularization on the cost function

is the new cost function convex?

$$\frac{1}{2N} \sum_{x,y \in \mathcal{D}} (y - w^\top x)^2 + \frac{\lambda}{2} \|w\|_2^2$$



Ridge regression

set the derivative to zero $J(w) = \frac{1}{2} \sum_{x,y \in \mathcal{D}} (y - w^\top x)^2 + \frac{\lambda}{2} w^\top w$

$$\nabla J(w) = \sum_{x,y \in \mathcal{D}} x(w^\top x - y) + \lambda w$$

$$= X^\top(Xw - y) + \lambda w = 0$$

linear system of equations $(X^\top X + \lambda \mathbf{I})w = X^\top y$

when using gradient descent, this term reduces the weights at each step (**weight decay**)

$$w = (X^\top X + \lambda \mathbf{I})^{-1} X^\top y$$


the only part different due to regularization

λI makes it invertible, adds a small value to the diagonals $X^\top X$

we can have linearly dependent features

the solution will be unique!




Example: polynomial bases



polynomial bases
 $\phi_k(x) = x^k$

Without regularization:





- using D=10 we can perfectly fit the data (high test error)



Example: polynomial bases

with regularization:

- fixed D=10, changing the amount of regularization



Probabilistic interpretation

recall linear regression & logistic regression maximize log-likelihood

$$w^{MLE} = \arg \max_w p(y|X, w)$$

linear regression $w^{MLE} = \arg \max_w \prod_{x,y \in \mathcal{D}} \mathcal{N}(y|w^\top x, \sigma^2)$

logistic regression $w^{MLE} = \arg \max_w \prod_{x,y \in \mathcal{D}} \text{Bernoulli}(y|\sigma(w^\top x))$

can we do Bayesian inference instead of maximum likelihood?

$$p(w|y, X) \propto p(w)p(y|w, X)$$

posterior

prior

likelihood

Maximum a Posteriori (MAP)

can we do Bayesian inference instead of maximum likelihood?

$$p(w|y, X) \propto p(w)p(y|w, X)$$

posterior prior likelihood

in general, this is expensive, but there's a cheap compromise:

$$\begin{aligned} \text{MAP estimate } w^{MAP} &= \arg \max_w p(w)p(y|X, w) \\ &= \arg \max_w \log p(y|X, w) + \log p(w) \end{aligned}$$

likelihood: original objective prior

all that is changing is the additional penalty on w

Gaussian Prior

$$\text{MAP estimate} \quad w^{MAP} = \arg \max_w \log p(y|X, w) + \underbrace{\log p(w)}_{\text{prior}}$$

assume independent zero-mean Gaussians

$$\mathcal{N}(\mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$


$$\log p(w) = \log \prod_{d=1}^D \mathcal{N}(w_d | 0, \tau^2) = - \sum_d \frac{w^2}{2\tau^2} + \text{const.}$$

does not depend on w
so it doesn't affect the optimization

lets call $\frac{1}{\tau^2} \rightarrow \lambda$

then we get the L2 regularization penalty $\frac{\lambda}{2} \|w\|_2^2$

smaller variance of the prior τ gives larger regularization λ



Laplace prior


another notable choice of prior is the Laplace distribution

minimizing negative log-likelihood $\rightarrow \sum_d \log p(w_d) = -\sum_d \frac{1}{\beta} |w_d| = -\frac{1}{\beta} ||w||_1$

L1 regularization: $J(w) \leftarrow J(w) + \lambda ||w||_1$ also called **lasso**

L1 norm of w

(least absolute shrinkage and selection operator)



$$p(w; \beta) = \frac{1}{2\beta} e^{-\frac{|w|}{\beta}}$$


notice the peak around zero

image from [here](#)

L_1 vs L_2 regularization

regularization path shows how $\{w_d\}$ change as we change λ

Lasso produces sparse weights (many are zero, rather than small)



red-line is the optimal λ from cross-validation, for lasso the model uses only 3 of the 8 features

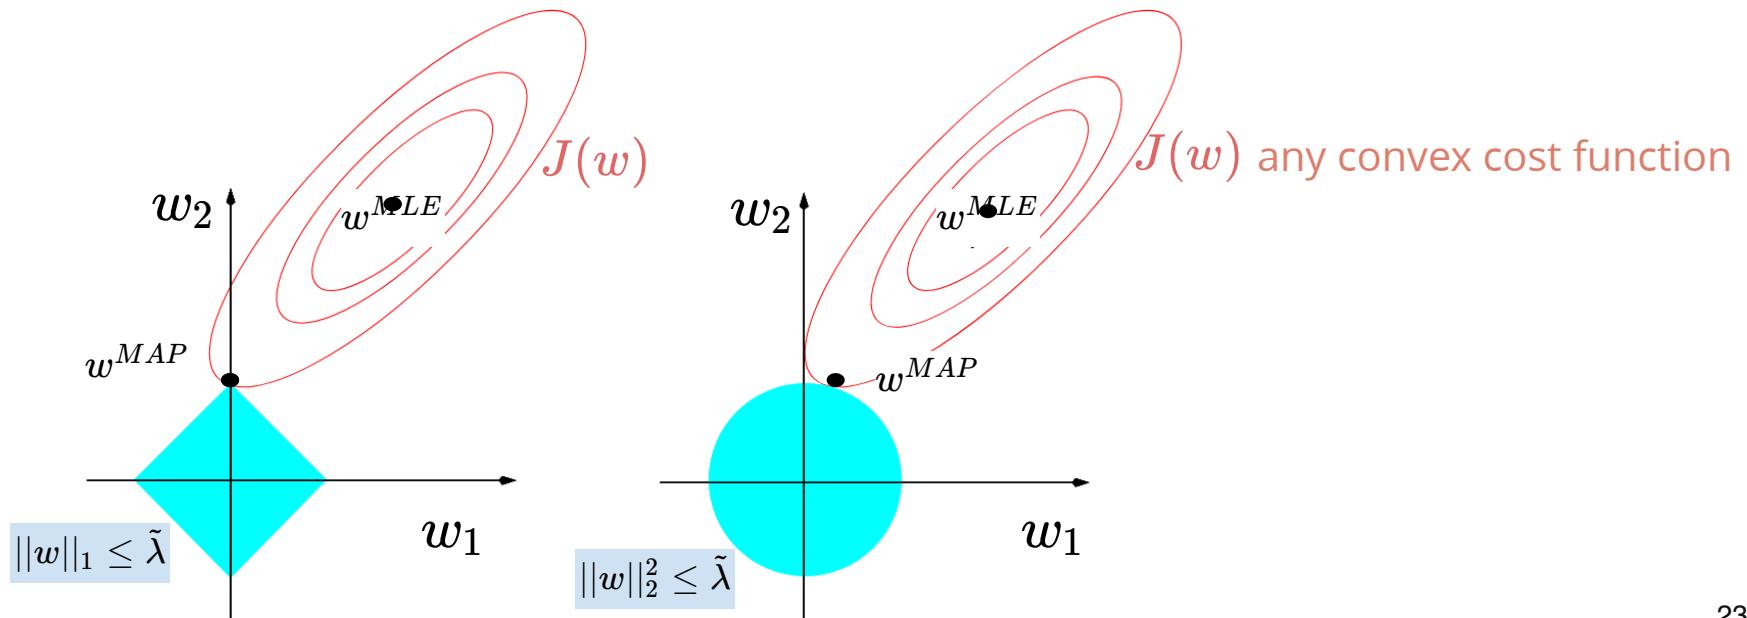
⇒ lasso results in sparse models

L_1 vs L_2 regularization

$$\min_w J(w) + \lambda ||w||_p^p$$


is equivalent to $\min_w J(w)$ subject to $||w||_p^p \leq \tilde{\lambda}$ for an appropriate choice of $\tilde{\lambda}$
 figures below show the constraint and the isocontours of $J(w)$

optimal solution with L1-regularization is more likely to have zero components



Subset selection

p-norms with $p \geq 1$ are convex (easier to optimize)




penalizes the **number of** features with non-zero weights

$$J(w) + \lambda ||w||_0 = J(w) + \lambda \sum_d \mathbb{I}(w_d \neq 0)$$

enforces a **penalty of λ** for each feature to be included in the model \Rightarrow performs feature selection

Subset selection

p -norms with $p \geq 1$ are convex (easier to optimize)



closer to **0-norm** → **L_0 norm**

L1 regularization is
a viable alternative
to L0 regularization

optimizing l_0 regularization
is a difficult *combinatorial*
problem: search over all 2^D
subsets


Adding L_2 regularization

do not penalize the bias w_0

L2 penalty makes the optimization easier too!

note that the optimal w_1 **shrinks**


example for **logistic regression**



```
1 def gradient(x, y, w, lambdaa):  
2     N,D = x.shape  
3     yh = logistic(np.dot(x, w))  
4     grad = np.dot(x.T, yh - y) / N  
5     grad[1:] += lambdaa * w[1:]  
6     return grad
```

weight decay

Sub-derivatives




L1 penalty is no longer smooth or differentiable (at 0)
extend the notion of derivative to non-smooth functions

sub-differential is the set of all **sub-derivatives** at a point

$$\partial f(\hat{w}) = \left[\lim_{w \rightarrow \hat{w}^-} \frac{f(w) - f(\hat{w})}{w - \hat{w}}, \lim_{w \rightarrow \hat{w}^+} \frac{f(w) - f(\hat{w})}{w - \hat{w}} \right]$$

if f is differentiable at \hat{w} then sub-differential has one member $\frac{d}{dw} f(\hat{w})$



another expression for sub-differential



$$\partial f(\hat{w}) = \{g \in \mathbb{R} \mid f(w) > f(\hat{w}) + g(w - \hat{w})\}$$

Subgradient

example

subdifferential for

$$f(w) = |w|$$



$$\partial f(0) = [-1, 1]$$

$$\partial f(w \neq 0) = \{\text{sign}(w)\}$$


recall, **gradient** was the vector of **partial derivatives**

subgradient is a vector of **sub-derivatives**

subdifferential for functions of multiple variables

$$\partial f(\hat{w}) = \{g \in \mathbb{R}^D \mid f(w) \geq f(\hat{w}) + g^\top (w - \hat{w})\}$$

we can use sub-gradient with diminishing step-size for optimization



Adding L_1 regularization




L1-regularized *linear regression* has efficient solvers

subgradient method for L1-regularized logistic regression

do not penalize the bias w_0

using diminishing learning rate

note that the optimal w_1 **becomes 0**



```
1 def gradient(x, y, w, lambdaa):
2     N,D = x.shape
3     yh = logistic(np.dot(x, w))
4     grad = np.dot(x.T, yh - y) / N
5     grad[1:] += lambdaa * np.sign(w[1:])
6     return grad
```

Regularization serves many purposes

$$w^* = (X^\top X)^{-1} X^\top y$$

$D \times 1$ $D \times N$ $N \times D$ $N \times 1$

what if $X^\top X$ is **not invertible**?

add a small value to the diagonals, a.k.a. **regularize**

what if **linear fit is not the best**?

use nonlinear basis

How to avoid **overfitting** then? **regularize**

what if **we want a sparse model**?

do feature selection and only keep important parameters with **regularizing**

Data normalization

what if we scale the input features, using different factors $\tilde{x}_d^{(n)} = \gamma_d x_d^{(n)} \forall d, n$

if we have **no regularization**: $\tilde{w}_d = \frac{1}{\gamma_d} w_d \forall d$

everything remains the same because: $\|Xw - y\|_2^2 = \|\tilde{X}\tilde{w} - y\|_2^2$

with regularization: $\|\tilde{w}\|_2 \neq \|w\|_2$ so the optimal **w** will be different!

features of different mean and variance will be penalized differently

normalization

$$\begin{cases} \mu_d = \frac{1}{N} x_d^{(n)} \\ \sigma_d^2 = \frac{1}{N-1} (x_d^{(n)} - \mu_d)^2 \end{cases}$$

makes sure all features have the same mean and variance $x_d^{(n)} \leftarrow \frac{x_d^{(n)} - \mu_d}{\sigma_d}$

we saw that this also helps with the optimization!

Summary

- complex models can overfit to training data
- regularization avoids this by penalizing model complexity
 - L1 & L2 regularization
 - probabilistic interpretation: different priors on weights
 - L1 produces sparse solutions (useful for feature selection)