

Applied Machine Learning

Multilayer Perceptron

Oumar Kaba



Exam post mortem

Learning objectives

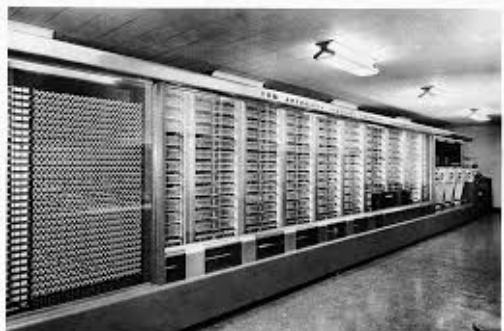
perceptron:

- model, objective, optimization

multilayer perceptron:

- model
 - different supervised learning tasks
 - activation functions
 - architecture of a neural network
- regularization techniques

Perceptron



old implementation (1960's)

historically a significant algorithm

(first neural network, or rather just a neuron)

biologically motivated model

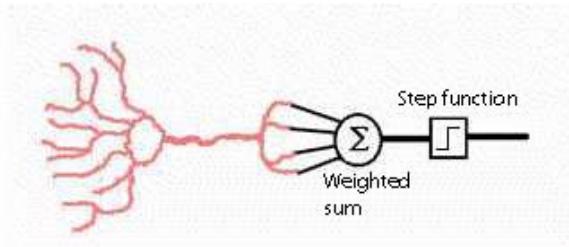
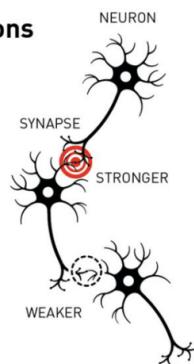


image:<https://cs.stanford.edu/people/eroberts/courses/soco/projects/neural-networks/Neuron/index.html>

Natural and artificial neurons

The brain's neural network is built from living cells, neurons, with advanced internal machinery. They can send signals to each other through the synapses. When we learn things, the connections between some neurons get stronger, while others get weaker.



The Nobel Prize in Physics 2024



Did you know that an artificial neural network is designed to mimic the brain?

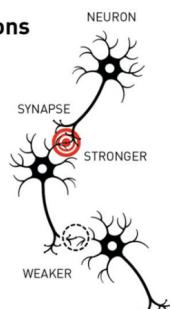
Inspired by biological neurons in the brain, artificial neural networks are large collections of “neurons”, or nodes, connected by “synapses”, or weighted couplings, which are trained to perform certain tasks. An artificial neural network processes information using its entire network structure. The inspiration initially came from the desire to understand how the brain works.

Learn more about this year's physics prize awarded for work on artificial neural networks: bit.ly/3Bi9H8u

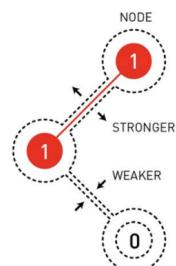
#NobelPrize

Natural and artificial neurons

The brain's neural network is built from living cells, neurons, with advanced internal machinery. They can send signals to each other through the synapses. When we learn things, the connections between some neurons gets stronger, while others get weaker.

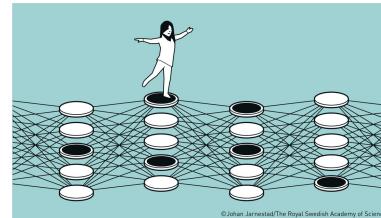


Artificial neural networks are built from nodes that are coded with a value. The nodes are connected to each other and, when the network is trained, the connections between nodes that are active at the same time get stronger, otherwise they get weaker.

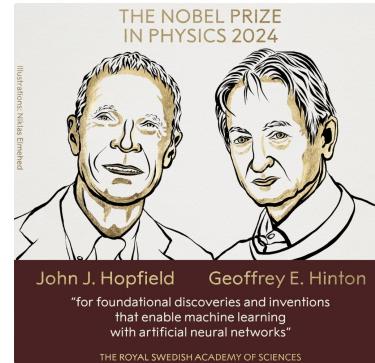


©Johan Jarnestad/The Royal Swedish Academy of Sciences

Last edited 6:19 AM · Oct 8, 2024 · 423.8K Views



© Johan Jarnestad/The Royal Swedish Academy of Sciences

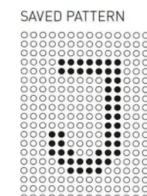
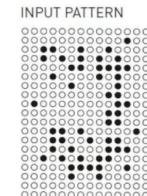
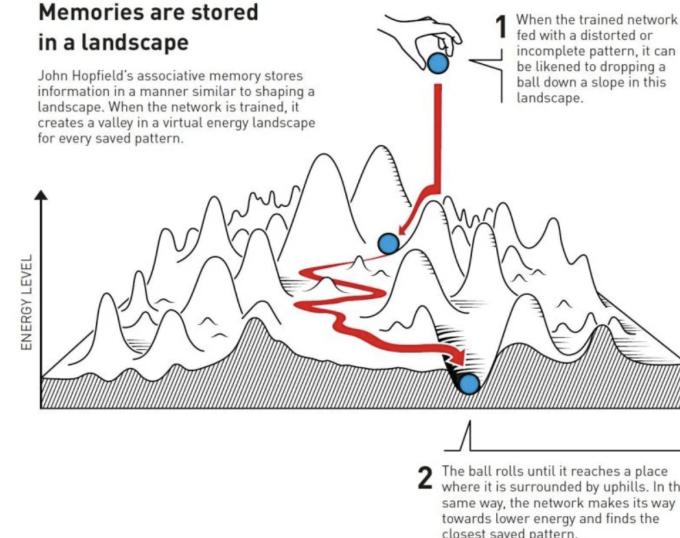


John J. Hopfield Geoffrey E. Hinton
“for foundational discoveries and inventions that enable machine learning with artificial neural networks”

THE ROYAL SWEDISH ACADEMY OF SCIENCES

Memories are stored in a landscape

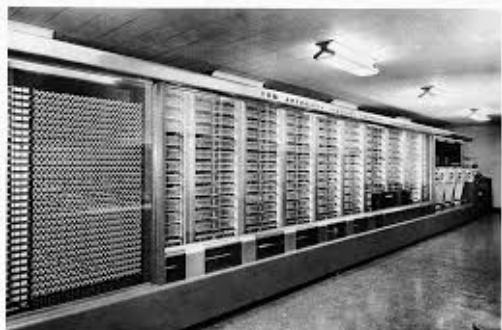
John Hopfield's associative memory stores information in a manner similar to shaping a landscape. When the network is trained, it creates a valley in a virtual energy landscape for every saved pattern.



© Johan Jarnestad/The Royal Swedish Academy of Sciences

read more [here](#)

Perceptron



old implementation (1960's)

historically a significant algorithm

(first neural network, or rather just a neuron)

biologically motivated model

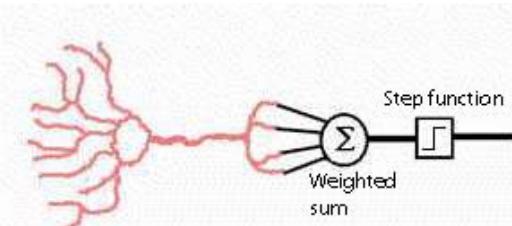
simple learning algorithm

convergence proof

beginning of *connectionist AI*

it's criticism in the book "Perceptrons" was a factor in AI winter

Model



$$f(x) = \text{sign}(w^\top x + w_0)$$

compare with models for linear and logistic regression:

$$f(x) = w^\top x + w_0$$

$$f(x) = \sigma(w^\top x + w_0)$$



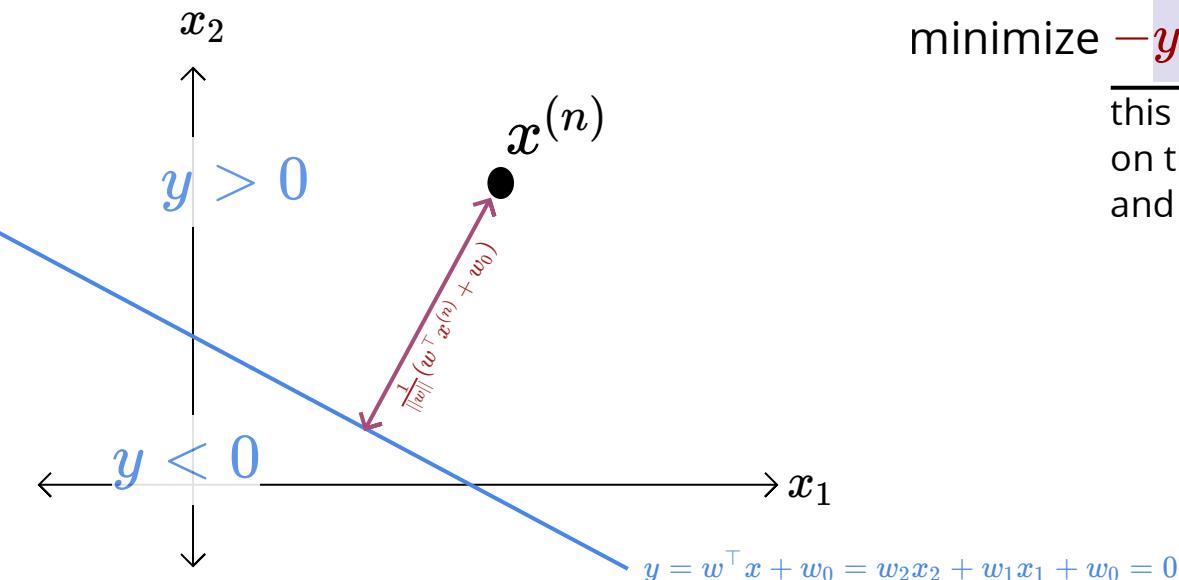
note that we're using +1/-1 for labels rather than 0/1.

Perceptron: objective

$$\hat{y}^{(n)} = \text{sign}(w^\top x^{(n)} + w_0)$$

misclassified if $y^{(n)}\hat{y}^{(n)} < 0$, try to make it positive

label and prediction have different signs



$$\leftarrow \hat{y}^{(n)} = \text{sign}(\downarrow) \rightarrow$$

$$\text{minimize } -y^{(n)}(w^\top x^{(n)} + w_0)$$

this is positive for points that are on the wrong side, minimize it and push them to the right side

Perceptron: optimization

if $y^{(n)}\hat{y}^{(n)} < 0$ minimize $J_n(w) = -y^{(n)}(w^\top x^{(n)})$
now we included bias in w

otherwise, do nothing

use stochastic gradient descent $\nabla J_n(w) = -y^{(n)}x^{(n)}$

$$w^{\{t+1\}} \leftarrow w^{\{t\}} - \alpha \nabla J_n(w) = w^{\{t\}} + \alpha y^{(n)} x^{(n)}$$

Perceptron uses learning rate of 1

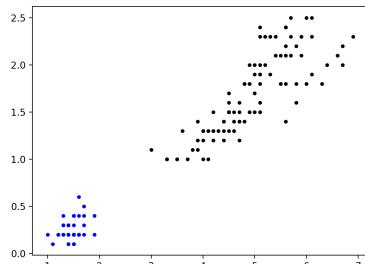
this is okay because scaling w does not affect prediction

$$\text{sign}(w^\top x) = \text{sign}(\alpha w^\top x)$$

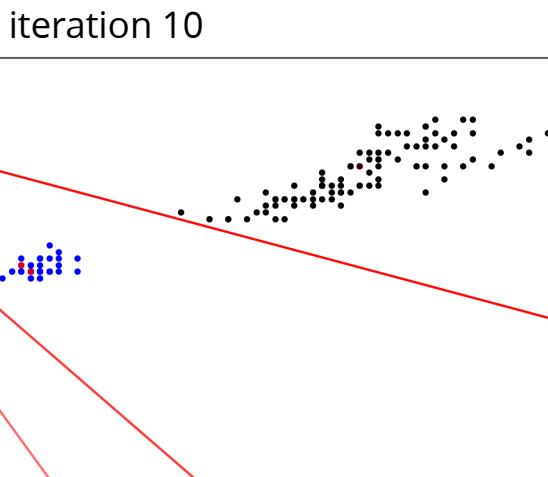
Perceptron convergence theorem

the algorithm is guaranteed to converge in finite steps if linearly separable

Perceptron: example



Iris dataset
(linearly separable case)



observations:

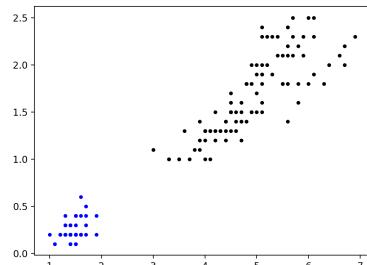
after finding a linear separator no further updates happen
the final boundary depends on the order of instances
(different from all previous methods)



```
1      N,D = x.shape
2      w = np.random.rand(D)
3      for t in range(max_iters):
4          n = np.random.randint(N)
5          yh = np.sign(np.dot(x[n,:], w))
6          if yh != y[n]:
7              w = w + y[n]*x[n,:]
```

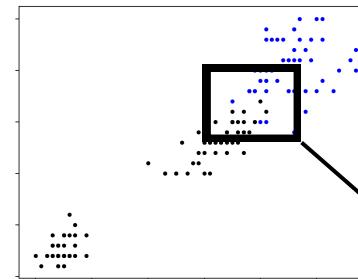
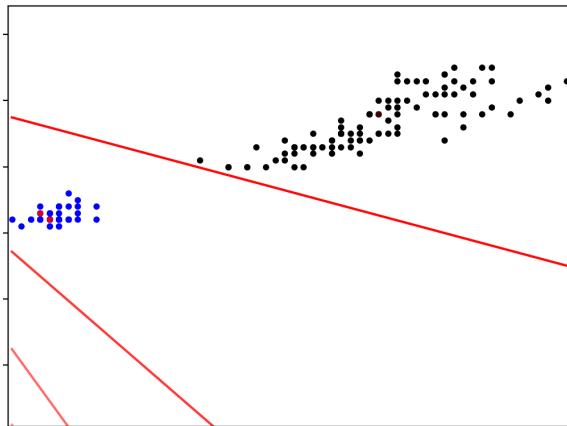
note that the code is not checking for convergence

Perceptron: example

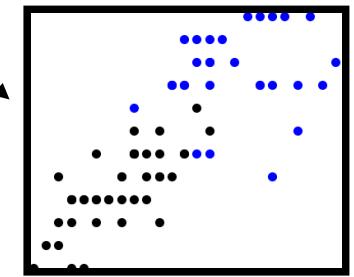


Iris dataset
(linearly separable case)

converged at iteration 10



Iris dataset
(**NOT** linearly
separable case)



the algorithm does
not converge
there is always a wrong
prediction and the weights
will be updated

Building more expressive model

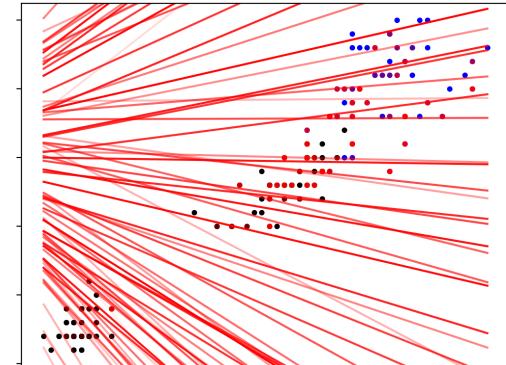
Perceptron is not expressive enough, can not model the data that is not linearly separable (gets stuck in cyclic updates)

how to increase the model's expressiveness?

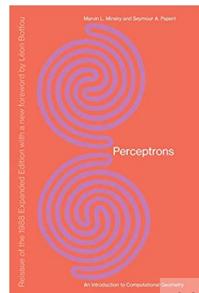
use **fixed nonlinear bases**: similar to what we have seen before

use **adaptive bases**: learn the parameters of the bases as well

- e.g., in regression $f(x) = \sum_m w_m \phi_m(x; \mathbf{v}_m)$



There is an influential book on the limitations of the perceptrons, see [here](#)



Adaptive Gaussian Bases

example

input has one dimension (D=1)

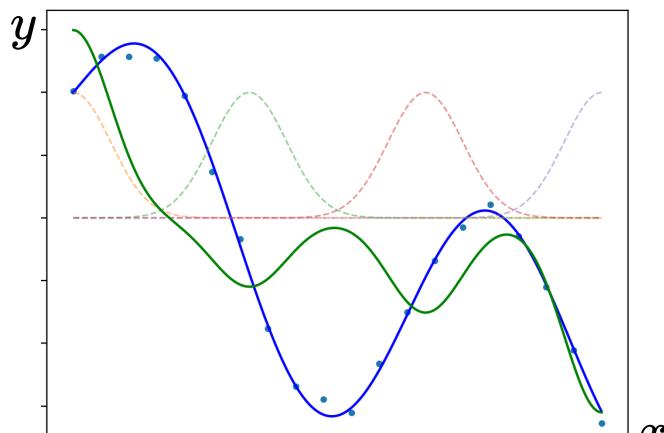
non-adaptive case

model: $f(x; w) = \sum_m w_m \phi_m(x)$

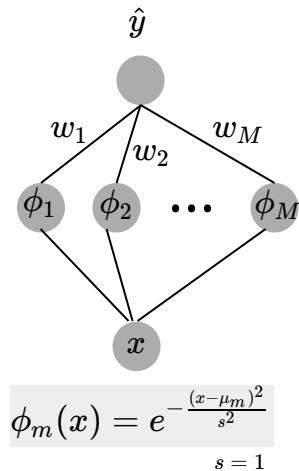
cost: $J(w) = \frac{1}{2} \sum_n (f(x^{(n)}; w) - y^{(n)})^2$

the model is linear in its parameters

the cost is convex in w



we have seen this before, centers (μ_m) are fixed



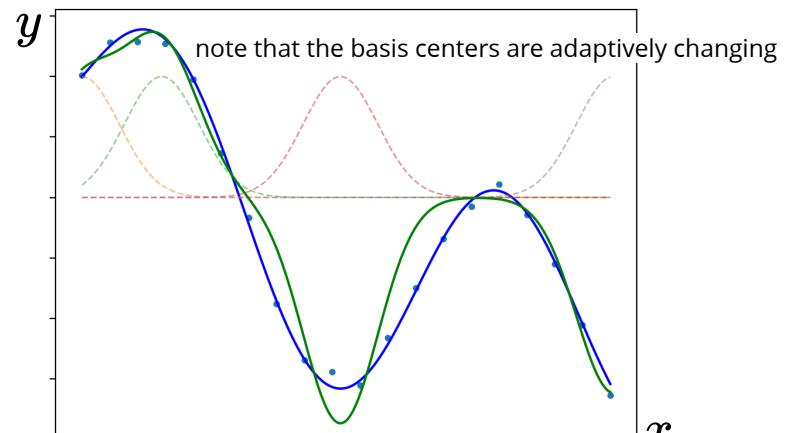
adaptive case

we can make the bases adaptive by learning the *centers*

model: $f(x; w, \mu) = \sum_m w_m \phi_m(x; \mu_m)$

not convex in all model parameters

use gradient descent to find a **local minimum**



adaptive case gives a better fit with the same number of bases (4)

Adaptive Sigmoid Bases

example input has one dimension (D=1)

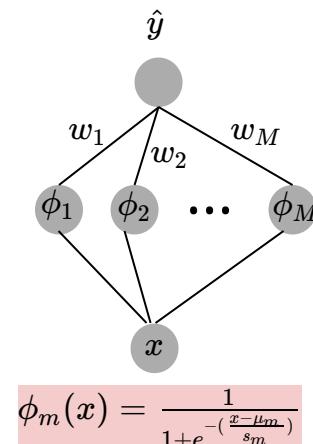
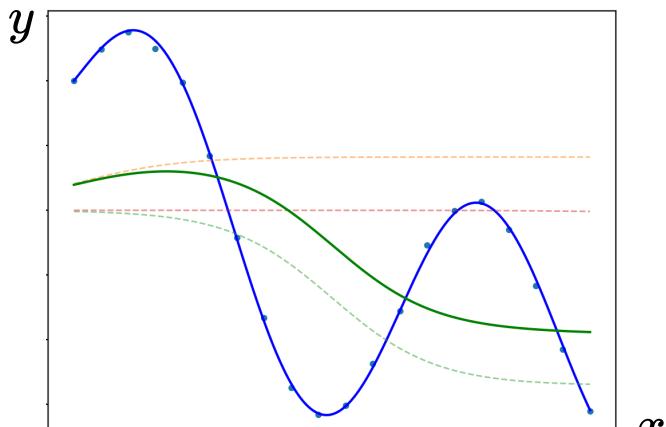
non-adaptive case

model: $f(x; w) = \sum_m w_m \phi_m(x)$

cost: $J(w) = \frac{1}{2} \sum_n (f(x^{(n)}; w) - y^{(n)})^2$

the model is linear in its parameters

the cost is convex in w



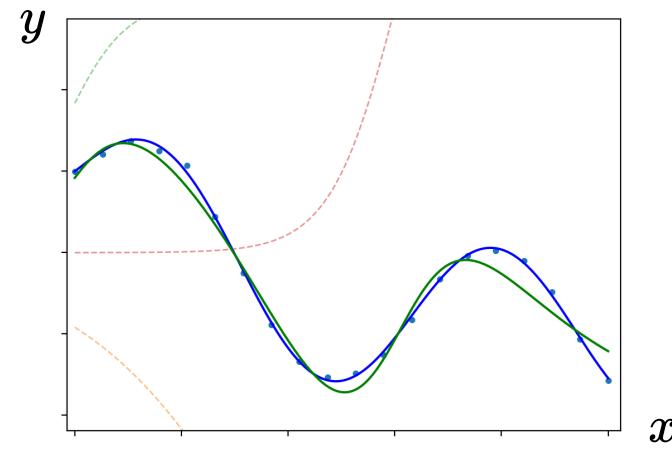
adaptive case

rewrite the sigmoid basis

$$\phi_m(x) = \sigma\left(\frac{x-\mu_m}{s_m}\right) = \sigma(v_m x + b_m)$$

model: $f(x; w, v, b) = \sum_m w_m \sigma(v_m x + b_m)$

optimize using gradient descent (find a local optima)



Adaptive Sigmoid Bases: General Case

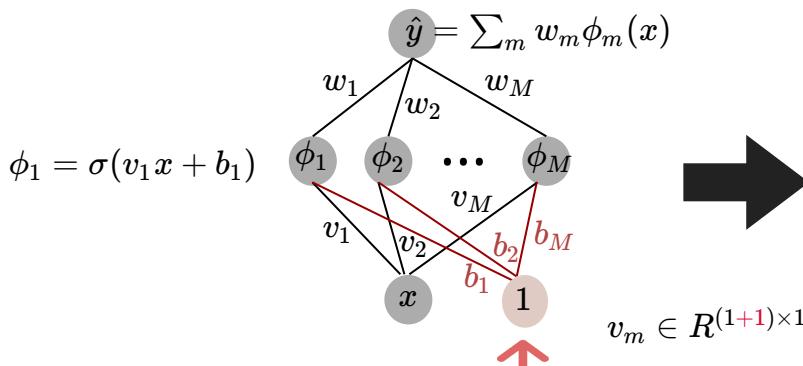
this is a **neural network** with two layers!!

each basis is the logistic regression model

$$\phi_m(x) = \sigma(v_m^\top x + b_m) \quad \forall m$$

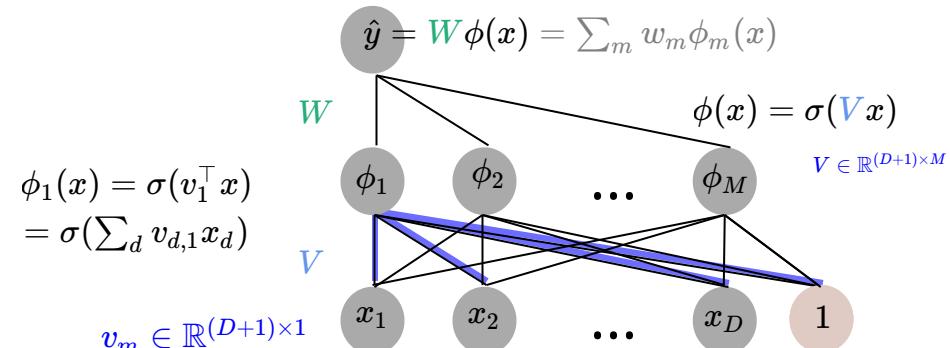
optimize V , W using gradient descent (find a local optima)

input has 1 dimension



this is the same as having a bias for each nonlinear basis

input has D dimension



Multilayer Perceptron (MLP)

suppose we have

- D inputs x_1, \dots, x_D
- C outputs $\hat{y}_1, \dots, \hat{y}_C$
- M hidden units z_1, \dots, z_M

model

$$\hat{y}_c = g \left(\sum_m W_{c,m} h \left(\sum_d V_{m,d} x_d \right) \right)$$

nonlinearity, activation function: we have different choices

more compressed form

$$\hat{y} = g(W h(Vx))$$

non-linearities are applied elementwise

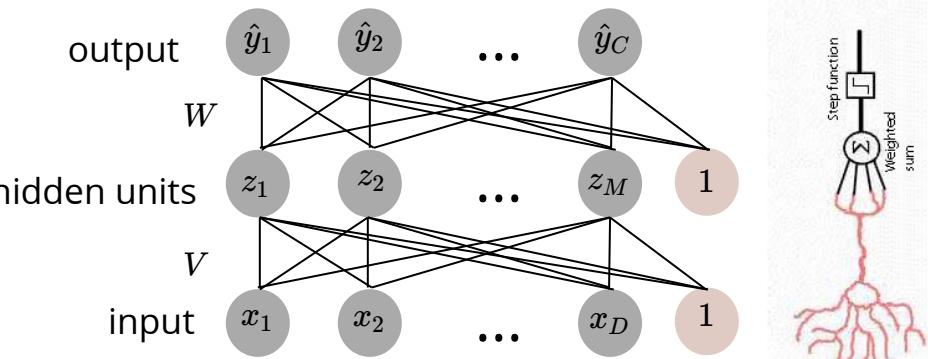
$$x \in \mathbb{R}^{D \times 1}$$

$$V \in \mathbb{R}^{M \times D}$$

$$Z = h(Vx) \in \mathbb{R}^{M \times 1}$$

$$W \in \mathbb{R}^{C \times M}$$

$$y \in \mathbb{R}^{C \times 1}$$



for simplicity we may drop bias terms

Regression using Neural Networks

the choice of **activation function** in the **final layer** depends on the task

model $\hat{y} = g(W h(V x))$

regression $\hat{y} = g(W z) = W z$

- we may have one or more output variables
- no activation (identity function)
- L2 loss = Gaussian likelihood

$$L(y, \hat{y}) = \frac{1}{2} \|y - \hat{y}\|_2^2 = -\log \mathcal{N}(y; \hat{y}, \mathbf{I}) + \text{constant}$$

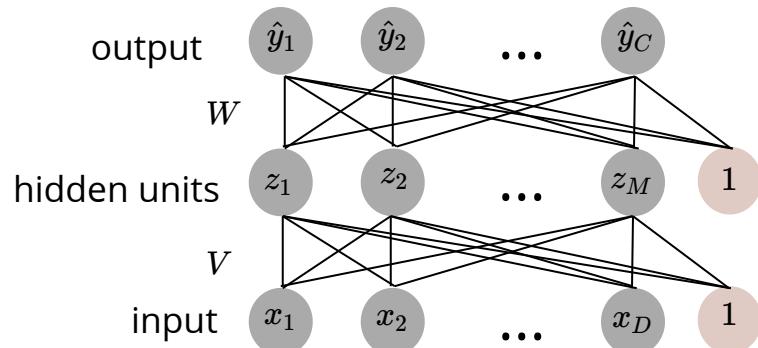
more generally

we may explicitly produce a distribution at output - e.g.,

- mean and variance of a Gaussian
- the loss will be the log-likelihood of the data under our model

$$L(y, \hat{y}) = \log p(y; f(x))$$

neural network outputs the parameters of a distribution



Classification using neural networks

the choice of activation function in the **final layer** depends on the task

model $\hat{y} = g(W h(V x))$

binary classification $\hat{y} = g(Wz) = \frac{1}{1+e^{-Wz}}$

- scalar output C=1
- activation function is logistic sigmoid
- CE loss = Bernoulli likelihood

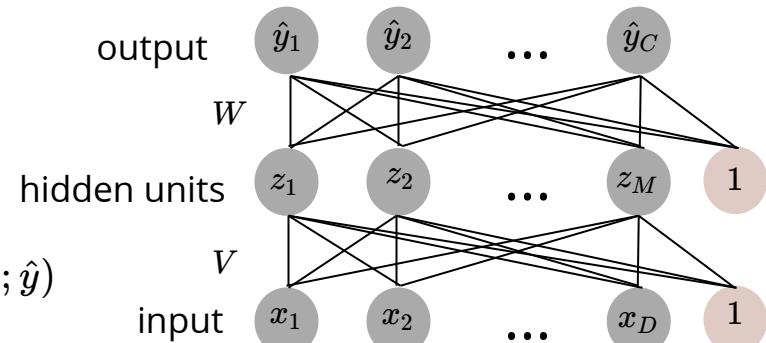
$$L(y, \hat{y}) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y}) = -\log \text{Bernoulli}(y; \hat{y})$$

multiclass classification $\hat{y} = g(Wz) = \text{softmax}(Wz)$

C is the number of classes

softmax activation

$$\text{multi-class cross entropy loss} = \text{categorical likelihood} \quad L(y, \hat{y}) = -\sum_k y_k \log \hat{y}_k = -\log \text{Categorical}(y; \hat{y})$$

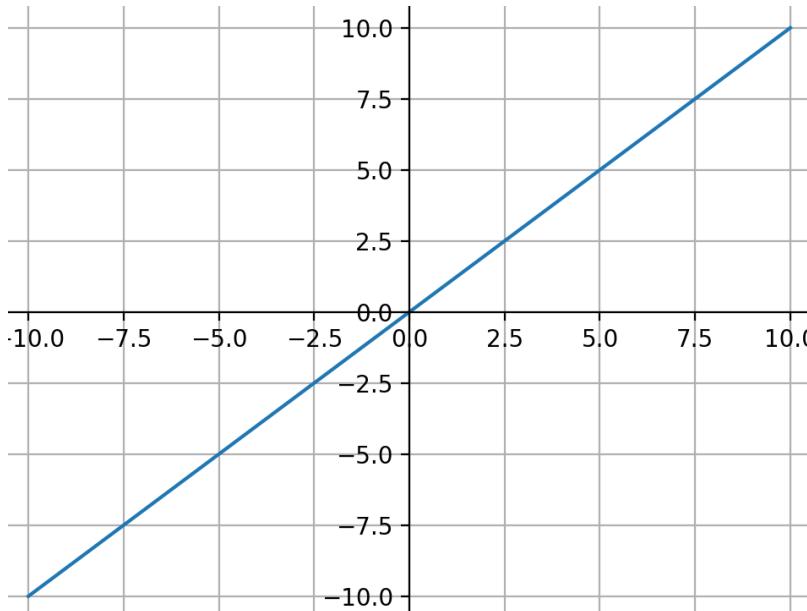


model

$$\hat{y} = g(W \mathbf{h}(V x))$$

Activation function

for **middle layer(s)** there is more freedom in the choice of activation function



$$h(x) = x \text{ identity (no activation function)}$$

composition of two linear functions is linear

$$\underbrace{WV}_{W'} x = W'x$$

$C \times M \quad M \times D \quad C \times D$

so nothing is gained (in representation power) by stacking linear layers

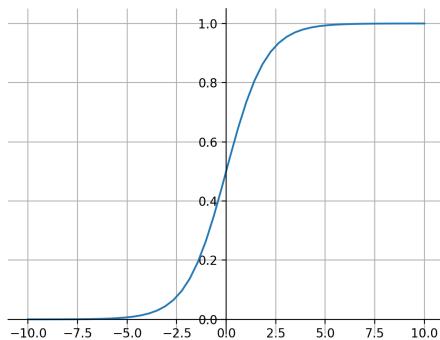
exception: if $M < \min(D, C)$ then the hidden layer is compressing the data (W' is low-rank)

model

$$\hat{y} = g(W \mathbf{h}(V x))$$

Activation function

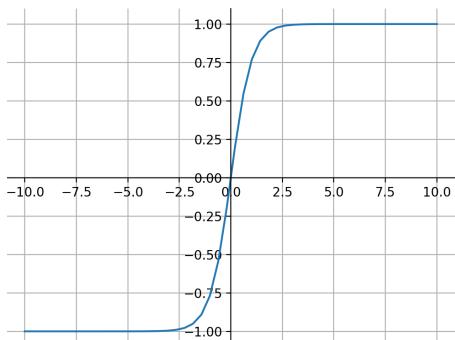
for **middle layer(s)** there is more freedom in the choice of activation function



$$h(x) = \sigma(x) = \frac{1}{1+e^{-x}} \text{ logistic function}$$

the same function used in logistic regression
used to be the function of choice in neural networks
away from zero it changes slowly, so the derivative is small (leads to vanishing gradient)
its derivative is easy to remember

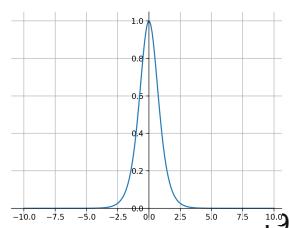
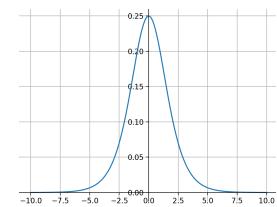
$$\frac{\partial}{\partial x} \sigma(x) = \sigma(x)(1 - \sigma(x))$$



$$h(x) = 2\sigma(x) - 1 = \frac{e^x - e^{-x}}{e^x + e^{-x}} \text{ hyperbolic tangent}$$

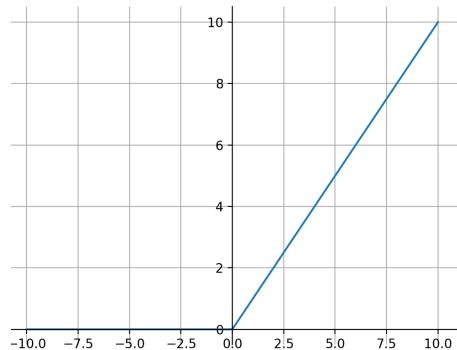
similar to sigmoid, but symmetric
often better for optimization because close to zero it
similar to a linear function (rather than an affine function when using logistic)
similar problem with vanishing gradient

$$\frac{\partial}{\partial x} \tanh(x) = 1 - \tanh(x)^2$$



Activation function

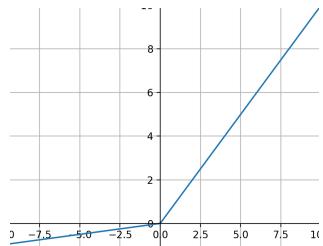
for **middle layer(s)** there is more freedom in the choice of activation function



$$h(x) = \max(0, x)$$
 Rectified Linear Unit (**ReLU**)

replacing logistic with ReLU significantly improves the training of deep networks
zero derivative if the unit is "inactive"
initialization should ensure active units at the beginning of optimization

leaky ReLU $h(x) = \max(0, x) + \gamma \min(0, x)$

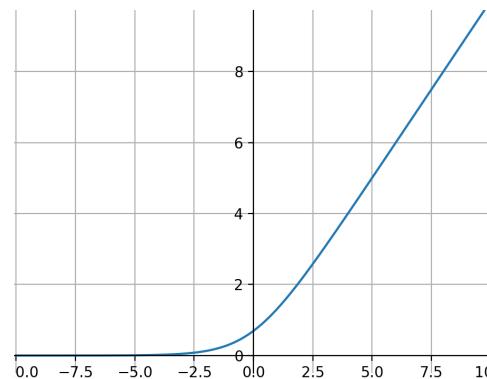


fixes the zero-gradient problem

parameteric ReLU:

make γ a learnable parameter

Softplus (differentiable everywhere)



$$h(x) = \log(1 + e^x)$$

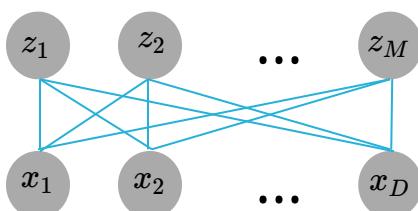
it doesn't perform as well in practice

Network architecture

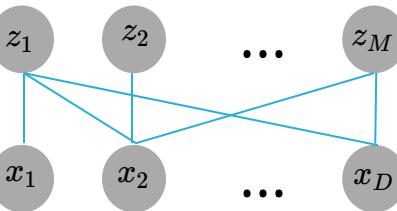
architecture is the overall structure of the network

feedforward network (aka multilayer perceptron)

- can have many layers
- # layers is called the **depth** of the network
- each layer can be **fully connected** (dense) or sparse

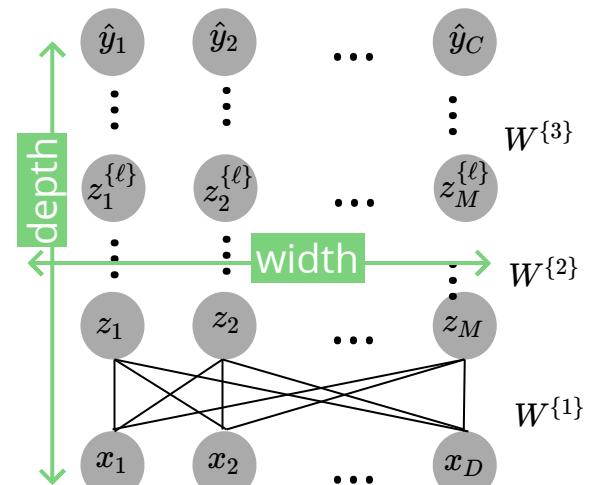


fully connected



sparsely connected

all outputs of one layer's units are input to
all the next units



$$z^{(l)} = h(W^{(l)} z^{(l-1)})$$

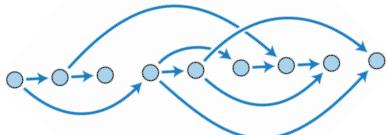
output of one layer is input to the next

Network architecture

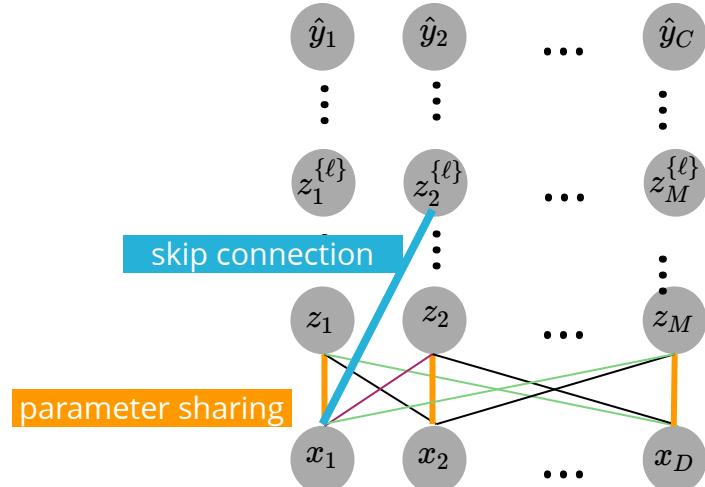
architecture is the overall structure of the network

feed-forward network (aka multilayer perceptron)

- can have many layers
- # layers is called the **depth** of the network
- each layer can be **fully connected** (dense) or sparse
- layers may have **skip layer connections**
- units may have different **activations**
- parameters may be shared across units (*e.g.*, in conv-nets)

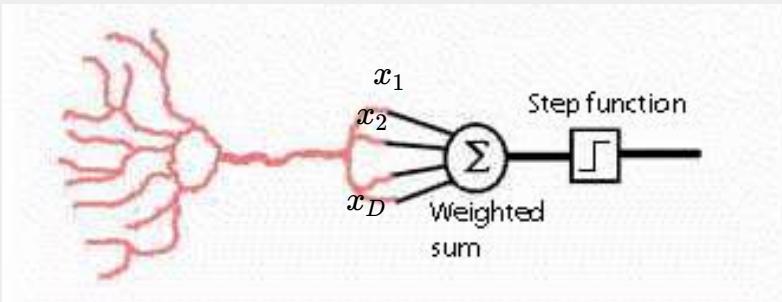


more generally a directed acyclic graph (DAG) expresses the feed-forward architecture

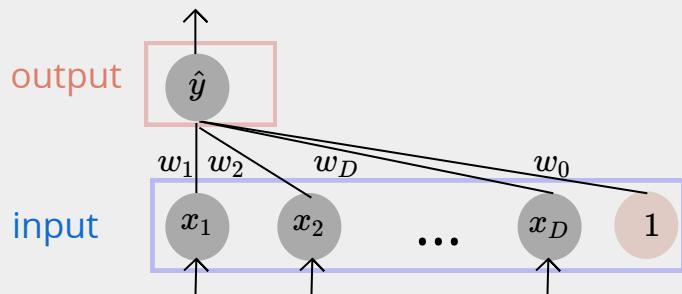


Multilayer Perceptron

Recall Perceptron



$$\hat{y} = \text{sign}\left(\sum_d w_d x_d + w_0\right)$$



$$\hat{y} = \text{sign}(w^\top x + w_0)$$

Example

$X = \begin{bmatrix} -x^{(1)\top} \\ -x^{(2)\top} \\ -x^{(3)\top} \\ -x^{(4)\top} \end{bmatrix} = \begin{bmatrix} 0 & 0 \\ 0 & 1 \\ 1 & 0 \\ 1 & 1 \end{bmatrix} y = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \end{bmatrix}$

$w_0 = 0 \quad w^\top x^{(i)}_{i \in [1..4]}** = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 2 \end{bmatrix}$

$w_0 = -1 \quad w^\top x - 1 = \begin{bmatrix} -1 \\ 0 \\ 0 \\ 1 \end{bmatrix}$

$\text{sign}^h(w^\top x) = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 1 \end{bmatrix} \quad \text{sign}^h(w^\top x - 1) = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$

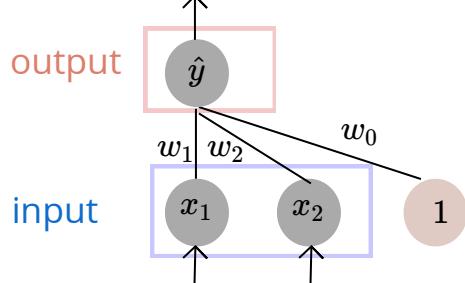
$$\text{sign}^h(x) = \mathbb{I}(x > 0)$$

Heaviside sign function, which is 0 for 0 and negative values

** we drop this for simplicity, it is similar to $X^\top W$, since $w^\top x$ is for one instance, however we use them interchangeably to show an affine function of input instances

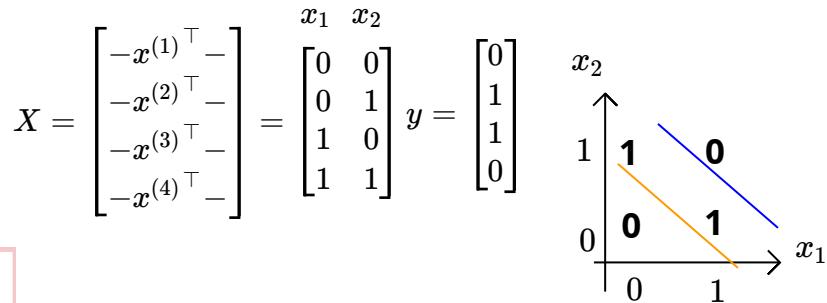
Both of these can not correctly classify this data

Multilayer Perceptron



$$w = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

$$\text{sign}^h(w^\top x) = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 1 \end{bmatrix}$$

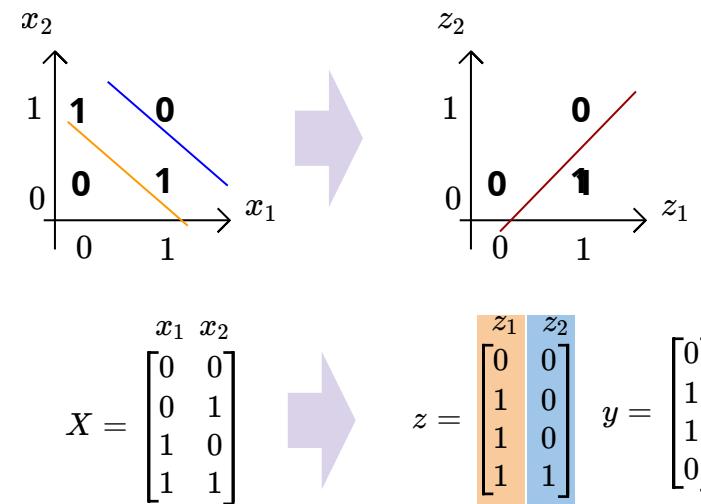


$$w_0 = 0 \quad w^\top x_{i \in [1..4]**} = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 2 \end{bmatrix} \quad w^\top x - 1 = \begin{bmatrix} -1 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$

$$\text{sign}^h(w^\top x - 1) = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$

Both of these can not correctly classify this data

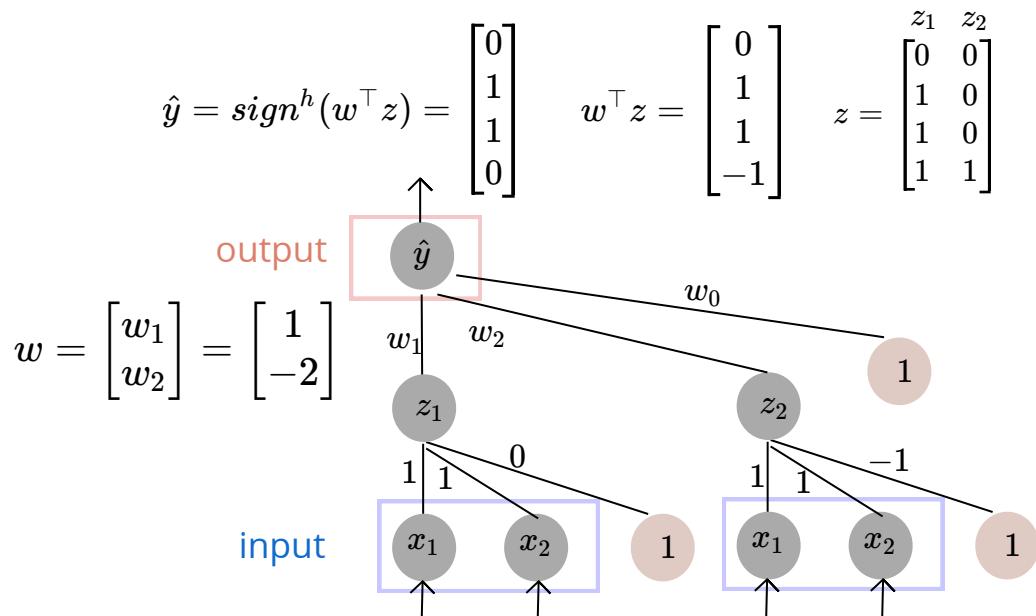
but together they transform the data so that it is linearly separable



Multilayer Perceptron

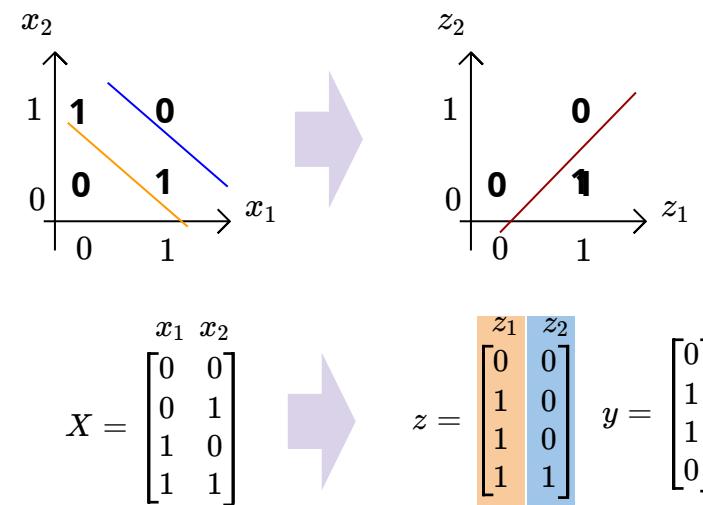
Example

put a linear classifier on it



Both of these can not correctly classify this data

but together they transform the data so that it is linearly separable



Multilayer Perceptron

$$V = \begin{bmatrix} 0 & 1 & 1 \\ -1 & 1 & 1 \end{bmatrix}$$

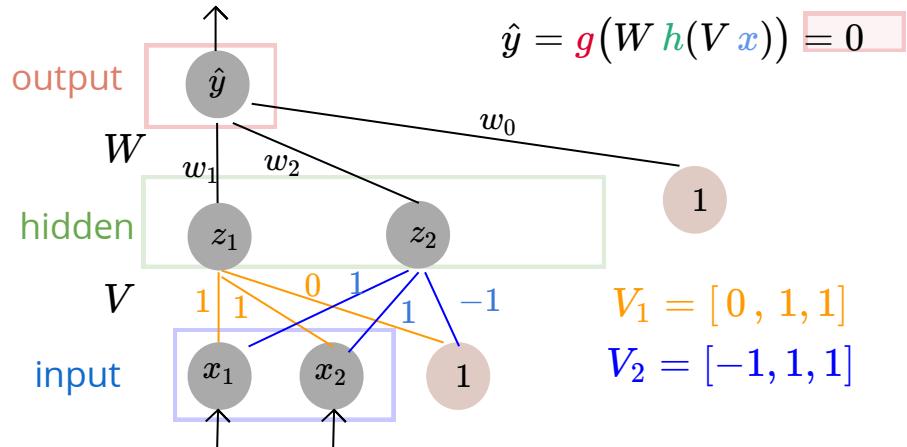
$$\underline{x} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}$$

$$V\underline{x} = \begin{bmatrix} 2 \\ 1 \end{bmatrix}$$

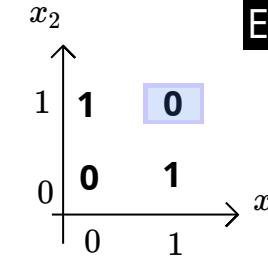
$$W = [0, 1, -2]$$

$$\hat{y} = g(W h(V \underline{x}))$$

$$\begin{aligned} h(V\underline{x}) &= \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} z_1 \\ W h(V\underline{x}) &= -1 \end{aligned}$$



$$\begin{bmatrix} -x^{(1)\top} \\ -x^{(2)\top} \\ -x^{(3)\top} \\ -x^{(4)\top} \end{bmatrix} = \begin{bmatrix} 0 & 0 \\ 0 & 1 \\ 1 & 0 \\ 1 & 1 \end{bmatrix} \quad y = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \end{bmatrix}$$



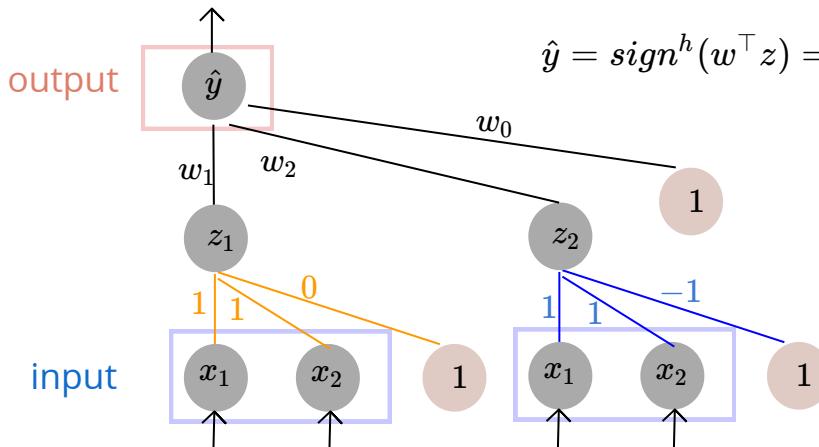
Example

$$z = \begin{bmatrix} z_1 & z_2 \\ 0 & 0 \\ 1 & 0 \\ 1 & 0 \\ 1 & 1 \end{bmatrix}$$

$$w = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} = \begin{bmatrix} 1 \\ -2 \end{bmatrix}$$

$$w^\top z = \begin{bmatrix} 0 \\ 1 \\ 1 \\ -1 \end{bmatrix}$$

$$\hat{y} = \text{sign}^h(w^\top z) = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \end{bmatrix}$$



Multilayer Perceptron

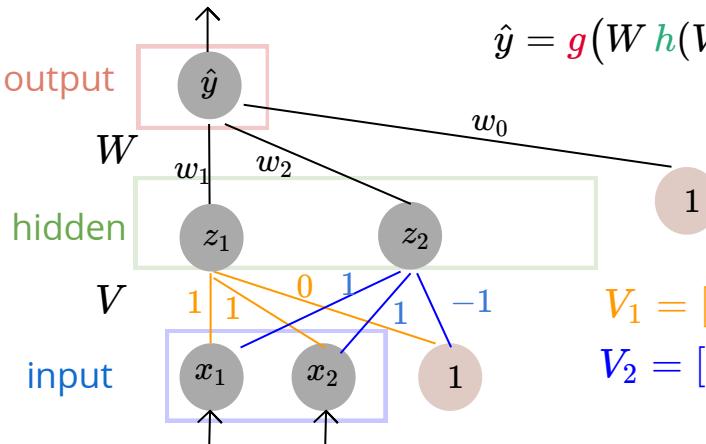
$$V = \begin{bmatrix} 0 & 1 & 1 \\ -1 & 1 & 1 \end{bmatrix}$$

$$\underline{x} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}$$

$$Vx = \begin{bmatrix} 2 \\ 1 \end{bmatrix}$$

$$W = [0, 1, -2]$$

$$\hat{y} = g(W h(V \underline{x}))$$



$$h(Vx) = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} z_1 \\ \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} z_2$$

$$W h(Vx) = -1$$

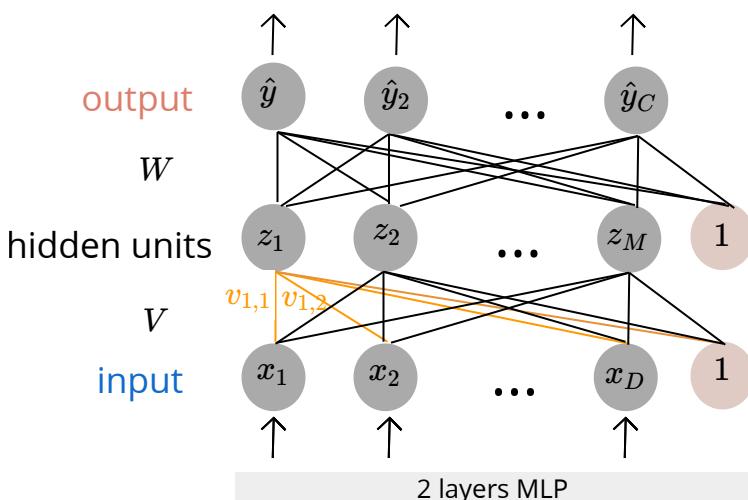
$$\hat{y} = g(W h(V x))$$

Example

$$V \in \mathbb{R}^{M \times \hat{D}} \quad W \in \mathbb{R}^{C \times \hat{M}}$$

$$z_m = h(V_m x) = h\left(\sum_d V_{m,d} x_d\right)$$

$$\hat{y}_k = g(W_k z) = g\left(\sum_m W_{k,m} z_m\right)$$

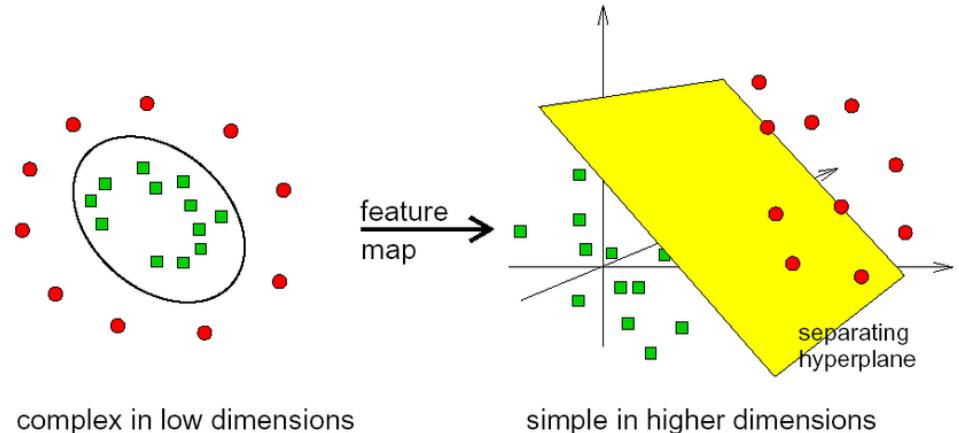
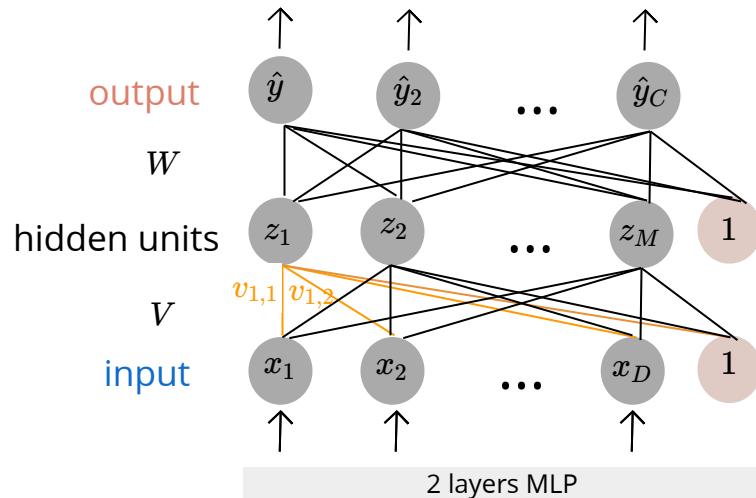


universal function approximator

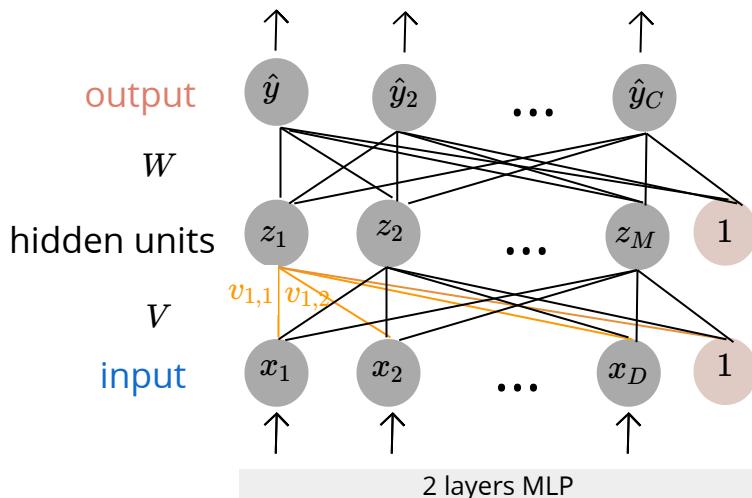
model any suitably smooth function, given enough hidden units, to any desired level of accuracy

Multilayer Perceptron

By increasing the number of hidden units,
we can increase expressivity



Multilayer Perceptron



universal function approximator

model any suitably smooth function, given enough hidden units, to any desired level of accuracy

Multilayer Feedforward Networks are Universal Approximators

KURT HORNIK

Technische Universität Wien

MAXWELL STINCHCOMBE AND HALBERT WHITE

University of California, San Diego

(Received 16 September 1988; revised and accepted 9 March 1989)

Abstract—This paper rigorously establishes that standard multilayer feedforward networks with as few as one hidden layer using arbitrary squashing functions are capable of approximating any Borel measurable function from one finite dimensional space to another to any desired degree of accuracy, provided sufficiently many hidden units are available. In this sense, multilayer feedforward networks are a class of universal approximators.

Keywords—Feedforward networks, Universal approximation, Mapping networks, Network representation capability, Stone-Weierstrass Theorem, Squashing functions, Sigma-Pi networks, Back-propagation networks.

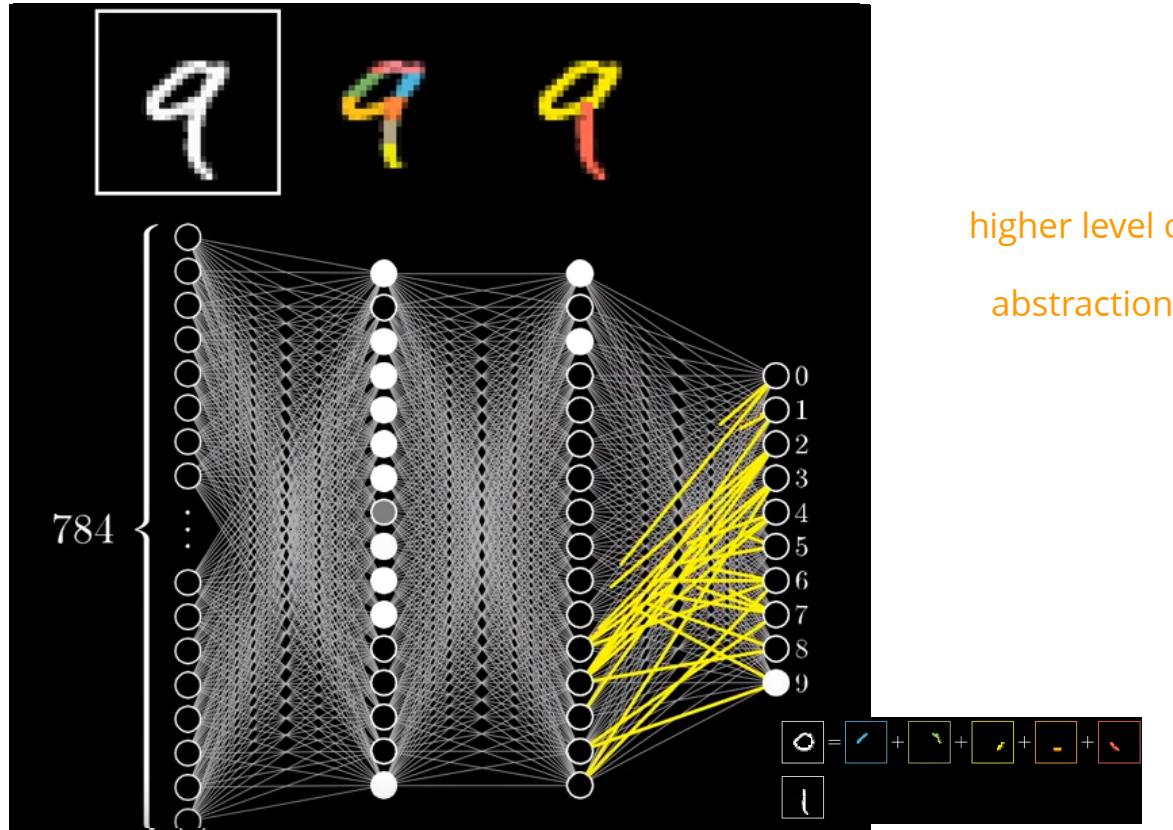
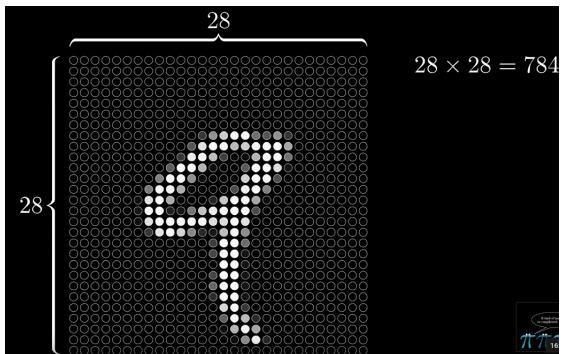
1. INTRODUCTION

It has been nearly twenty years since [Minsky and Papert \(1969\)](#) conclusively demonstrated that the

any function encountered in applications leads one to wonder about the ultimate capabilities of such networks. Are the successes observed to date re-

MNIST Example

classifying handwritten digits



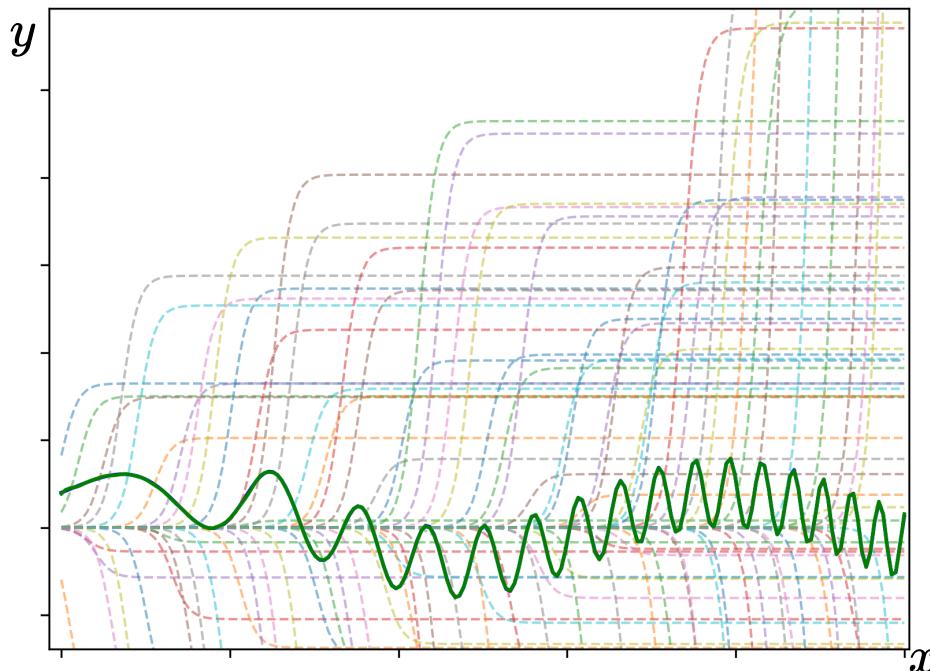
see this video for better intuition

https://www.youtube.com/watch?v=aircAruvnKk&list=PLZHQBObOWTQDNU6R1_67000Dx_ZCJB-3pi&index=2&t=7s

Expressive power

universal approximation theorem

an MLP with single hidden layer can approximate any continuous function with arbitrary accuracy



for 1D input we can see this even with **fixed bases**
 $M = 100$ in this example

the fit is good (hard to see the blue line)

however # bases (M) should grow exponentially
with D (**curse of dimensionality**)

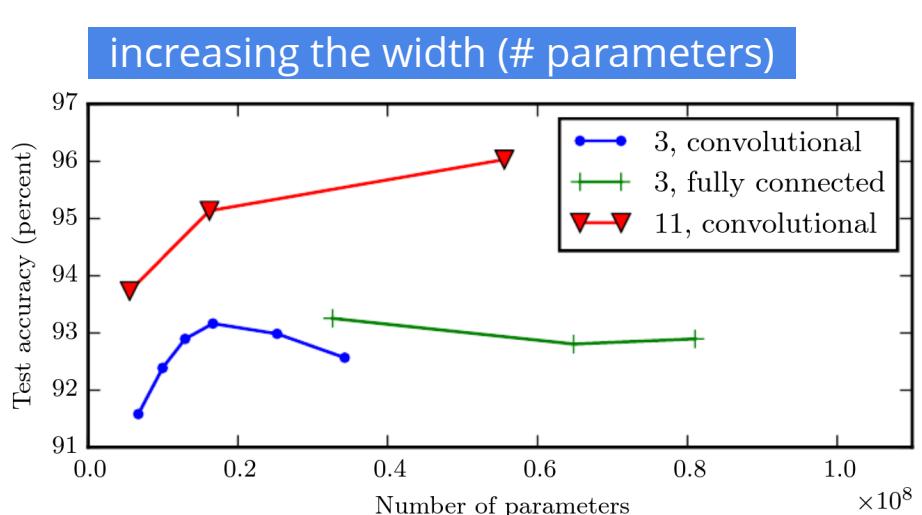
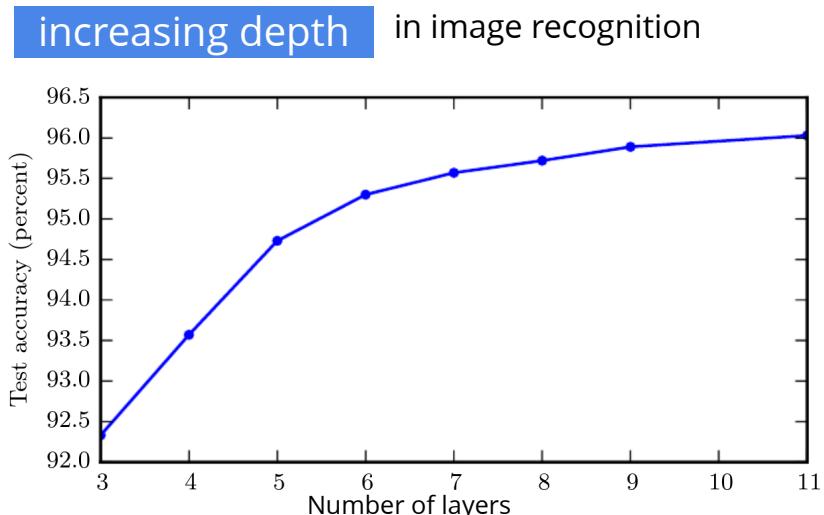
Caveats of the universality

- we may need a very wide network (large M)
- this is only about training error, we care about test error

Depth vs Width

Deep networks (with ReLU activation) of bounded width are also shown to be universal

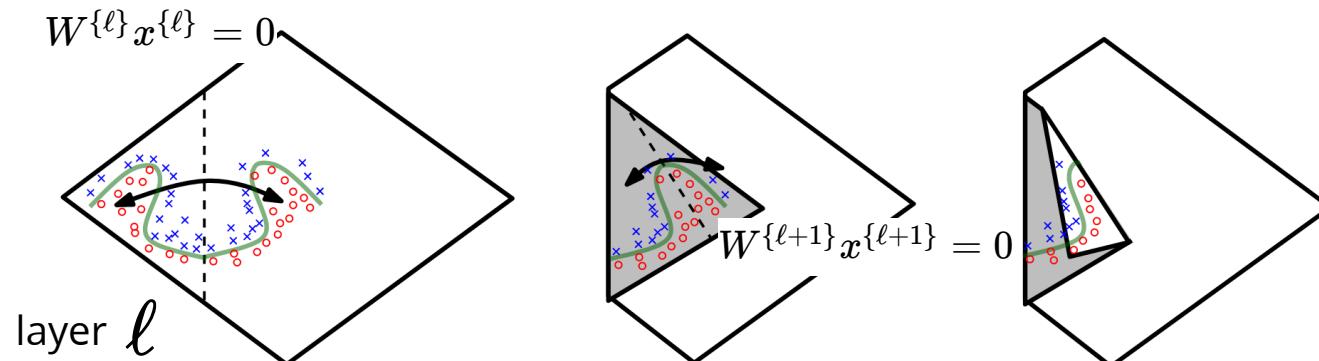
- empirically, increasing the depth is often more effective than increasing the width (#parameters per layer)
- compositional functional form through depth is a useful inductive bias



Depth vs Width

Deep networks (with ReLU activation) of bounded width are also shown to be universal
number of regions (in which the network is linear) grows exponentially with depth

simplified demonstration $h(W^{\{\ell\}}x) = |W^{\{\ell\}}x|$



Regularization strategies

universality of neural networks also means they can overfit
strategies for variance reduction:

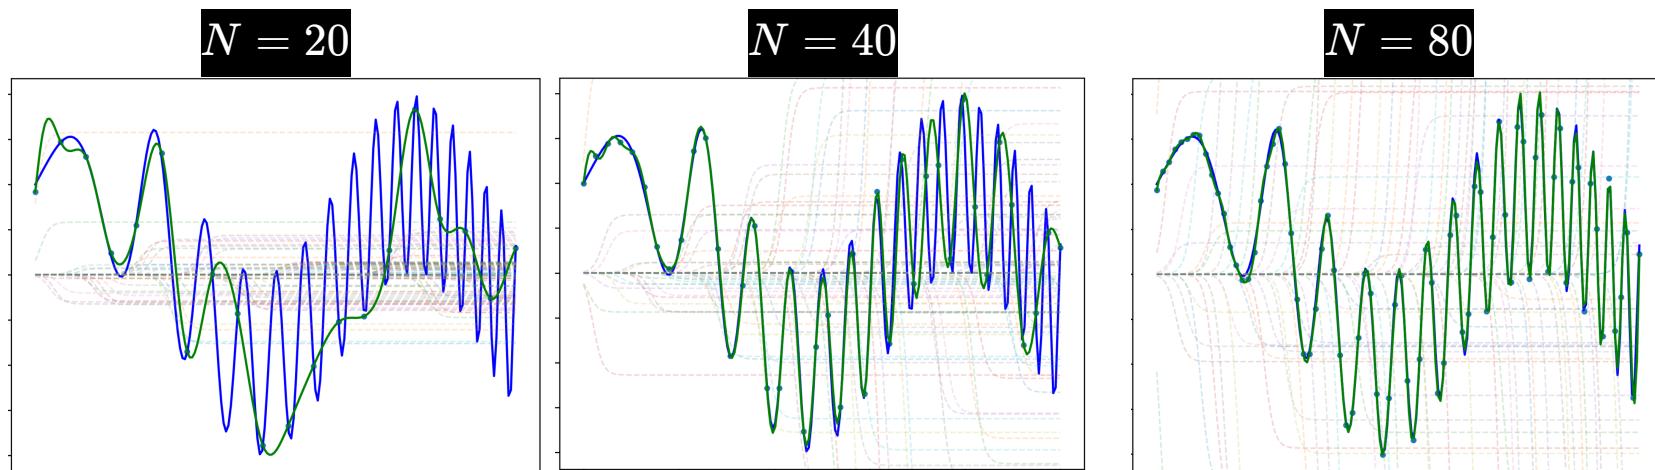
- L1 and L2 regularization (*weight decay*)
- data augmentation
- noise robustness
- early stopping
- dropout
- bagging
- sparse representations (e.g., *L1 penalty on hidden unit activations*)
- semi-supervised and multi-task learning
- adversarial training
- parameter-tying

Regularization using Data augmentation

a larger dataset results in a better generalization

example: in all 3 examples below training error is close to zero

however, a larger training dataset leads to better generalization



Regularization using Data augmentation

a larger dataset results in a better generalization



idea



increase the size of dataset by adding reasonable transformations $\tau(x)$ that change the label in predictable ways; e.g., $f(\tau(x)) = f(x)$

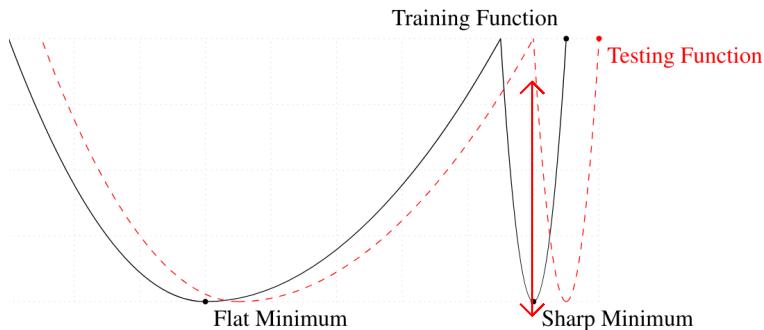
special approaches to data-augmentation

- adding noise to the input
- adding noise to hidden units
 - noise in higher level of abstraction
- learn a **generative model** $\hat{p}(x, y)$ of the data
 - use $x^{(n')}, y^{(n')} \sim \hat{p}$ for training

sometimes we can achieve the same goal by designing the models that are **invariant** to a given set of transformations

Regularization using Noise robustness

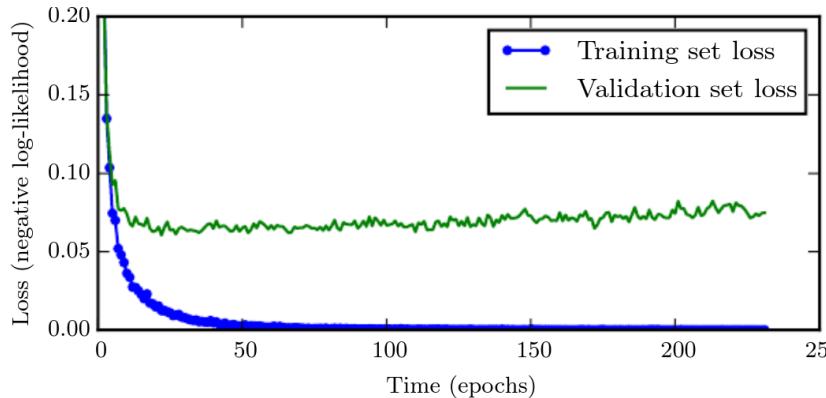
1. **input** (data augmentation)
2. **hidden units** (e.g., in dropout as we see soon)
3. **weights** the cost is not sensitive to small changes in the weight (**flat minima**)



flat minima generalize better
good performance of SGD using small minibatch is attributed to converging to flat minima which generalizes better (train loss closer to test loss)
in this case, SGD regularizes the model due to **gradient noise**

<https://arxiv.org/pdf/1609.04836.pdf>

Regularization using Early stopping



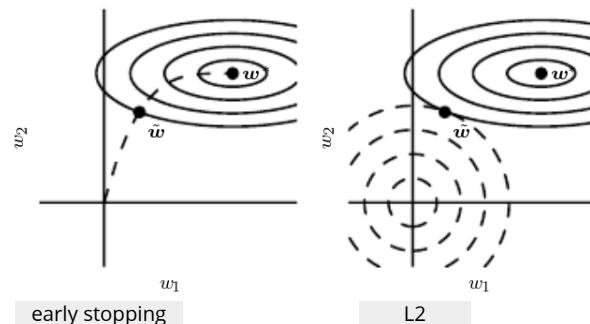
the **test loss-vs-time step** is "often" U-shaped
use validation for early stopping
also saves computation!

early stopping bounds the region of the parameter-space that is reachable in T time-steps

assuming

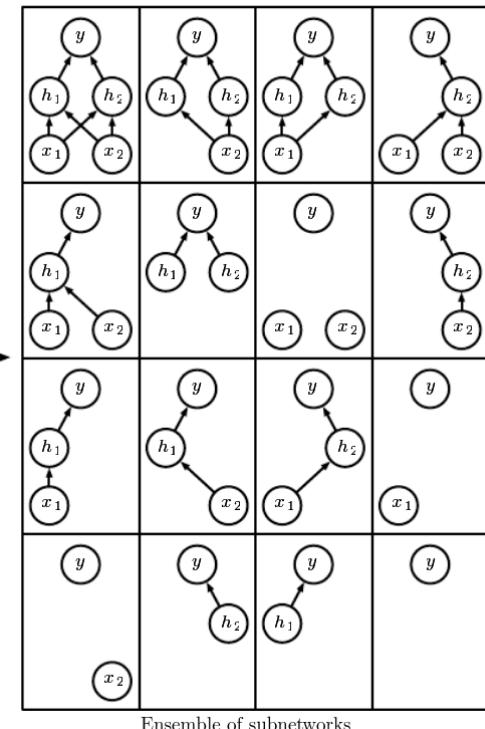
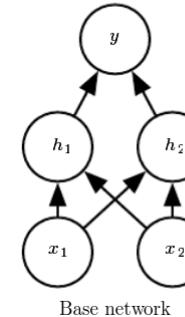
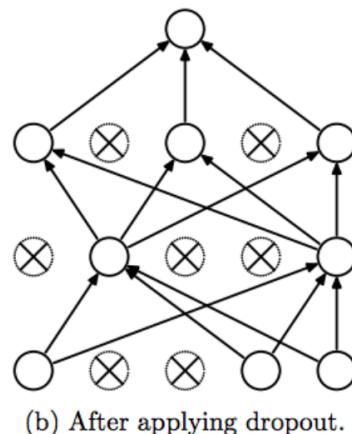
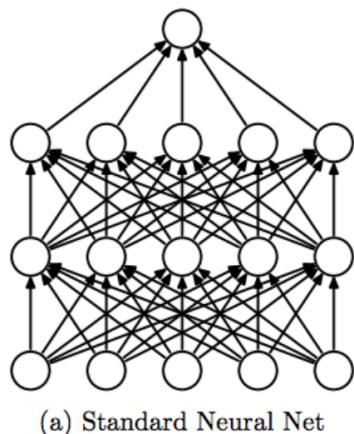
- *bounded gradient*
- *starting with a small w*

it has an effect similar to L2 regularization
we get the regularization path (various λ)



Regularization using Dropout

randomly remove a subset of units during training



can be viewed as exponentially many subnetworks that share parameters
is one of the most effective regularization schemes for MLPs

Regularization using Dropout

during training

for each instance (n):

randomly dropout each unit with probability p (e.g., $p=.5$)

only the remaining subnetwork participates in training

at test time

ideally we want to average over the prediction of all possible sub-networks

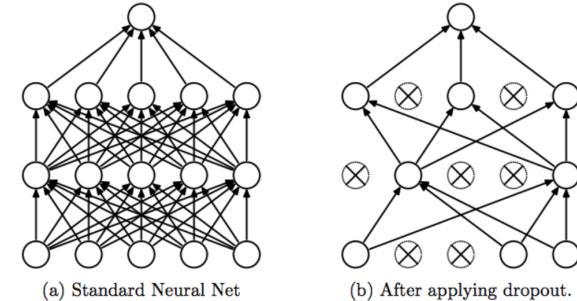
this is computationally infeasible, instead:

1) Monte Carlo dropout: average the prediction of several feed-forward passes using dropout

2) weight scaling: scale the weights by p to compensate for dropout

e.g., for 50% dropout, scale by a factor of 2

either multiply by 2 in training or divide by 2 at the end of training



Interactive neural network training

Useful platform to visually experiment with training small networks

<https://playground.tensorflow.org/>

Summary

Deep feed-forward networks learn **adaptive bases**
more complex bases at higher layers
increasing **depth** is often preferable to width
various choices of **activation function** and **architecture**
universal approximation power
their expressive power often necessitates using **regularization** schemes